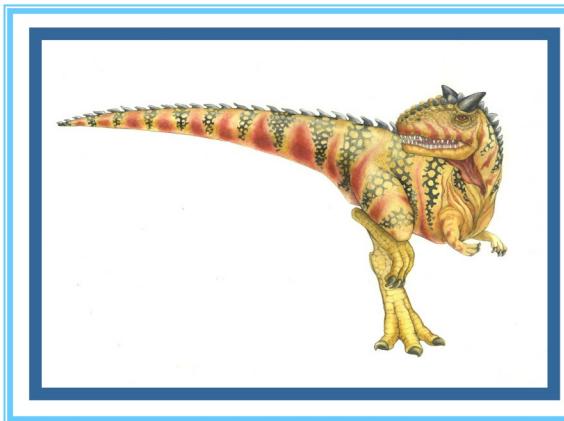


Chapter 6: CPU Scheduling





Chapter 6: CPU Scheduling

- Basic Concepts
- Scheduling Criteria
- Scheduling Algorithms
- Examples





Objectives

Multi ^{prog}
processor

- To introduce CPU scheduling, which is the basis for multiprogrammed operating systems
- To describe various CPU-scheduling algorithms
- To discuss evaluation criteria for selecting a CPU-scheduling algorithm for a particular system





Basic Concepts

- Maximum CPU utilization obtained with multiprogramming (MPg)
- Continuous Cycle :
 - ✓ one process has to wait (I/O)
 - ✗ Operating system takes the CPU away
 - ✓ Give CPU to another process
 - This pattern continues
- CPU–I/O Burst Cycle – Process execution consists of a cycle of CPU execution and I/O wait

✓ (C) P_0 - executing
✗ (C) P_0 - wait I/O
✓ (C) P_1 - executing





CPU and I/O Burst Cycle

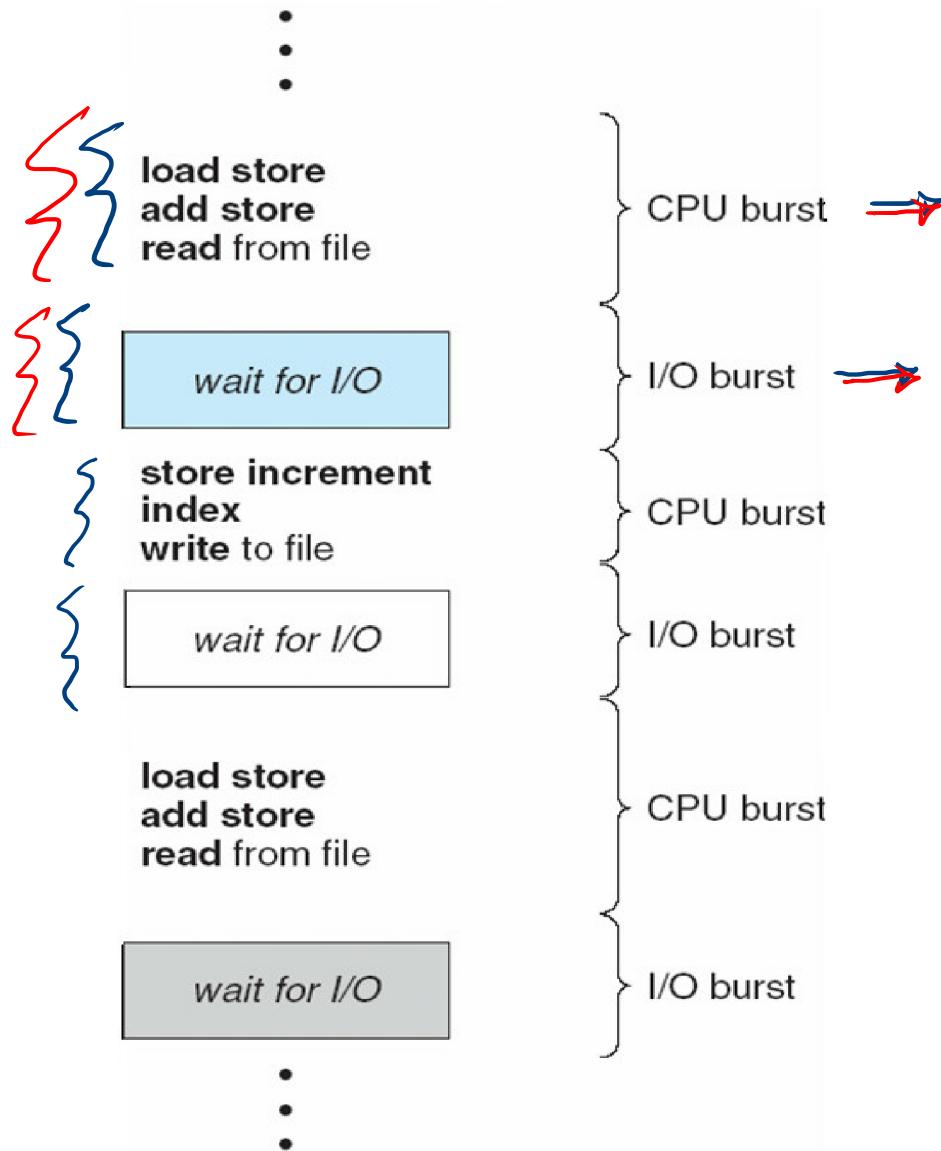
- Almost all processes alternate between two states in a continuing **cycle**, as shown in Figure below :
 - A CPU burst of performing calculations, and
 - An I/O burst, waiting for data transfer in or out of the system.

- Processes alternate back and forth between this two states.





Alternating Sequence of CPU and I/O Bursts





CPU Scheduler

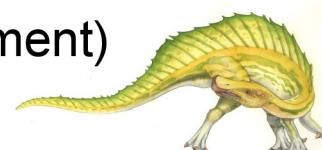
- Selects from among the processes in ready queue, and allocates the CPU to one of them
 - FIFO queue
 - Priority queue
 - Tree
 - Unordered linked-list
- CPU scheduling decisions may take place when a process:
 1. Switches from running to waiting state (I/O request)
 2. Switches from running to ready state (e.g. when interrupt occurs)
 3. Switches from waiting to ready (e.g. at completion of I/O)
 4. Terminates
- Scheduling under 1 and 4 is **nonpreemptive**
- All other scheduling is **preemptive**
 - Consider access to shared data
 - Consider preemption while in kernel mode
 - Consider interrupts occurring during crucial OS activities





Scheduling Criteria

- ✓ **CPU utilization** – keep the CPU as busy as possible
- ✓ **Throughput** – # of processes that complete their execution per time unit
- ✓ **Turnaround time**
 - amount of time to execute a particular process
 - the interval from the time of submission of a process to the time of the completion.
 - sum of the periods spent waiting to get into memory, waiting in the ready queue, executing on the CPU, doing I/O
$$(\tau_c - \tau_s)$$
- ✓ **Waiting time** – amount of time a process has been waiting in the ready queue
- ✓ **Response time** – amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)





Scheduling Algorithm Optimization Criteria

- Max CPU utilization
- Max throughput
- Min turnaround time
- Min waiting time
- Min response time





First-Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- Suppose that the processes arrive in the order: P_1, P_2, P_3
The Gantt Chart for the schedule is:

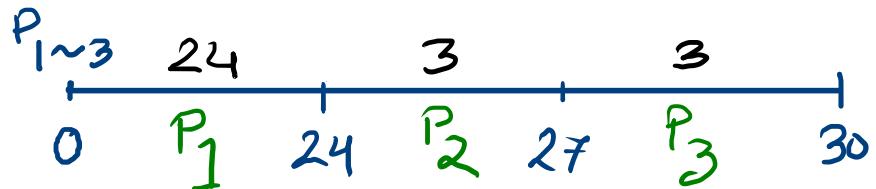


- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: $(0 + 24 + 27)/3 = 17$



Section - 1

P_1	24
P_2	3
P_3	3
$P_1 \rightarrow P_2 \rightarrow P_3$	



* Turnaround time

$$\left. \begin{array}{l} P_1: 24 - 0 = 24 \\ P_2: 27 - 0 = 27 \\ P_3: 30 - 0 = 30 \end{array} \right\} \text{Avg} = 21$$

* Waiting time

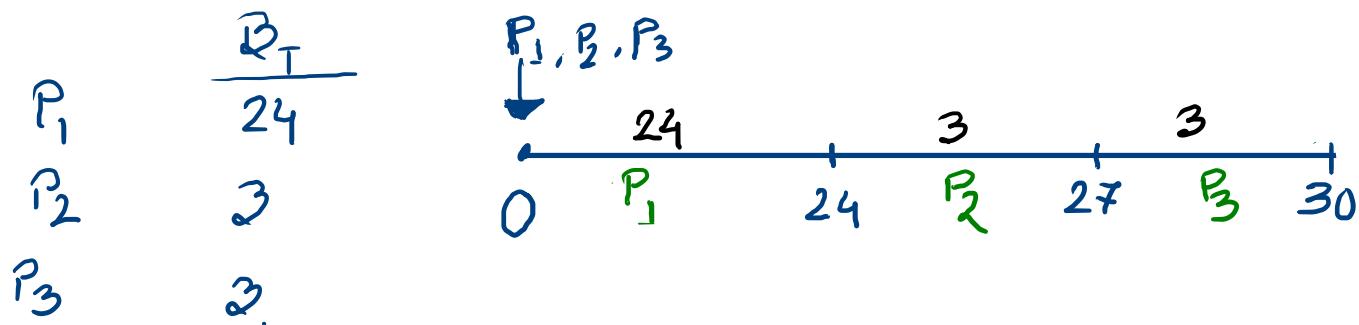
$$\left. \begin{array}{l} P_1: 0 - 0 = 0 \\ P_2: 24 - 0 = 24 \\ P_3: 27 - 0 = 27 \end{array} \right\} \text{Avg} = 17$$

* Response time

$$\left. \begin{array}{l} P_1: 0 - 0 = 0 \\ P_2: 24 - 0 = 24 \\ P_3: 27 - 0 = 27 \end{array} \right\} \text{Avg} = 17$$

$$\underbrace{(C_t - A_t)}_{\downarrow} + \boxed{\quad}$$

Section - 02



④ $P_1 \rightarrow P_2 \rightarrow P_3$ ④ Turnaround times

$$\left. \begin{array}{l} P_1 : 24 - 0 = 24 \\ P_2 : 27 - 0 = 27 \\ P_3 : 30 - 0 = 30 \end{array} \right\} \text{Avg} \rightarrow T_A = 27$$

→ ④ Waiting times.

$$\left. \begin{array}{l} P_1 : 0 \\ P_2 : 24 \\ P_3 : 27 \end{array} \right\} \text{Avg} \rightarrow W_t = 17$$

④ Non-preemptive
 W_t, R_t converge

④ Response times

$$\left. \begin{array}{l} P_1 : 0 \\ P_2 : 24 \\ P_3 : 27 \end{array} \right\} \text{Avg} \rightarrow R_t = 17$$

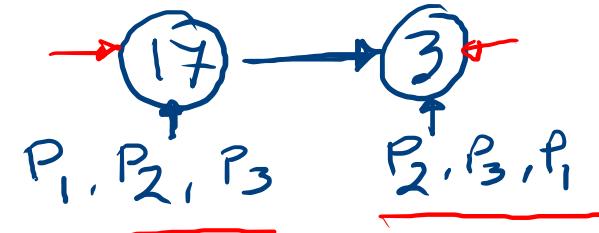
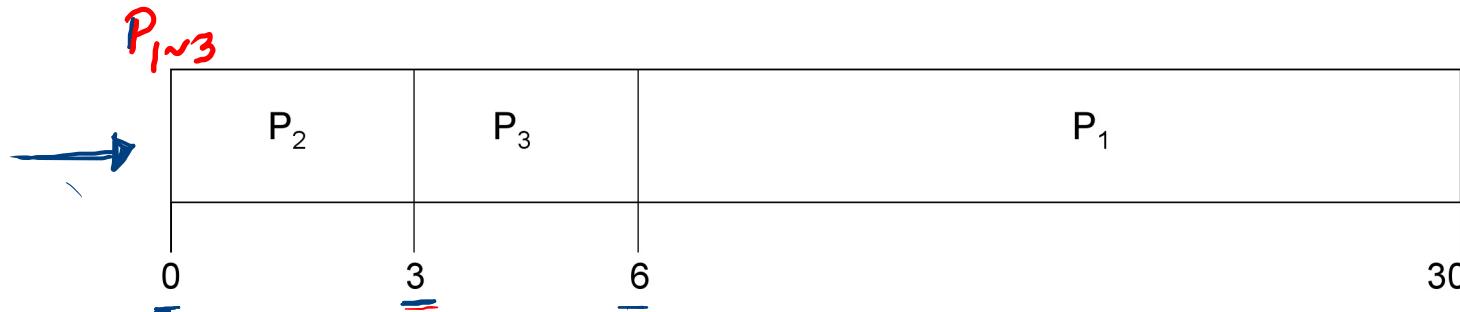


FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order:

$$\rightarrow \underline{P_2, P_3, P_1}$$

- The Gantt chart for the schedule is:



Waiting time for P₁ = 6; P₂ = 0; P₃ = 3

Average waiting time: $(6 + 0 + 3)/3 = 3$

Much better than previous case

* **Convoy effect** - short process behind long process

Consider one CPU-bound and many I/O-bound processes





FCFS Scheduling (Cont.)

Convoy Effect :

many I/O bound process and one CPU bound process



CPU bound process	I/O bound process	Effect
I/O device	I/O queue	CPU site idle
CPU processing	Ready queue	I/O site idle





Shortest-Job-First (SJF) Scheduling

- ~~Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time~~
- Two schemes:
 - ~~Non-preemptive~~ – once CPU is given to the process it cannot be preempted until it completes its CPU burst
 - ~~preemptive~~ – if a new process arrives with CPU burst length less than remaining time of current executing process, preempt. This scheme is known as the Shortest-Remaining-Time-First (SRTF)
- SJF is optimal – gives minimum average waiting time for a given set of processes

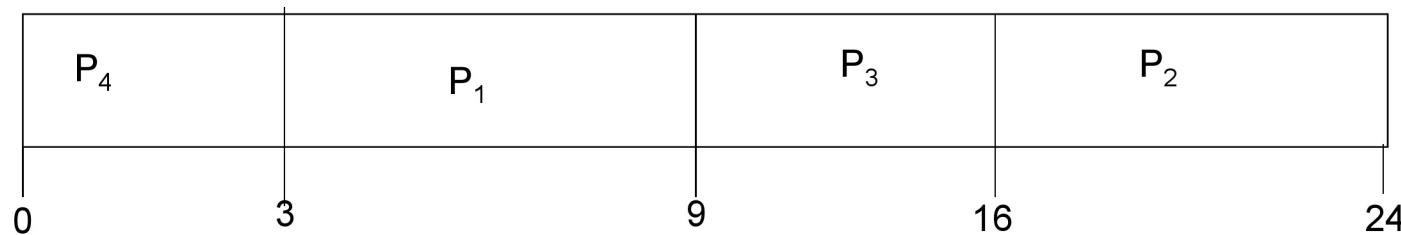




Example of SJF

<u>Process</u>	<u>Burst Time</u>
P_1	6
P_2	8
P_3	7
P_4	3

- SJF scheduling chart



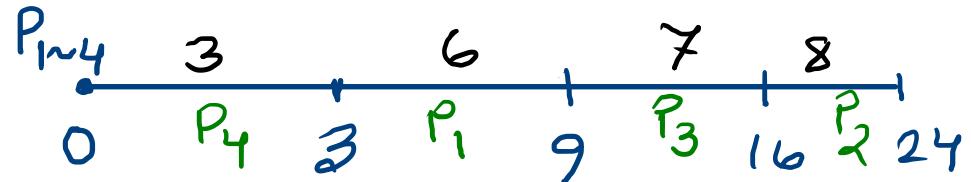
$$\text{Average waiting time} = (3 + 16 + 9 + 0) / 4 = 7$$



Section -1

P ₁	6
P ₂	8
P ₃	7
P ₄	3

P₄ → P₁ → P₃ → P₂



Waiting times.

$$\left. \begin{array}{l} P_1: 3-0 = 3 \\ P_2: 16-0 = 16 \\ P_3: 9-0 = 9 \\ P_4: 0-0 = 0 \end{array} \right\} \text{Avg: } 7$$

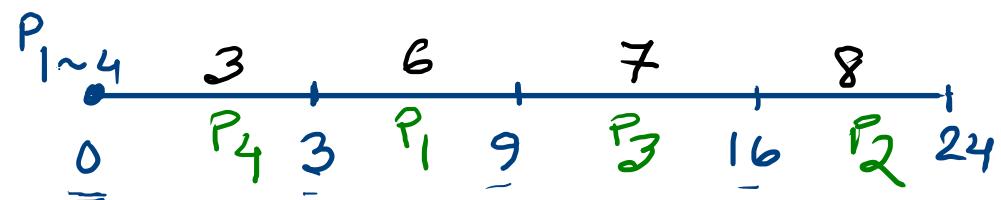
SJF = Sort + FCF

$SJF = \text{Sort followed by FCFS}$

Section - 02

$P_1 \quad C$
 $P_2 \quad 8$
 $P_3 \quad 7$
 $P_4 \quad 3$

$\rightarrow P_4 - P_1 - P_3 - P_2$



* Waiting Times $\rightarrow (C_t - A_t) + \boxed{\quad}$ $(Q_{Ct} - A_t)$

$P_1 \rightarrow 3 - 0 = 3$
 $P_2 \rightarrow 16 - 0 = 16$
 $P_3 \rightarrow 9 - 0 = 9$
 $P_4 \rightarrow 0 - 0 = 0$

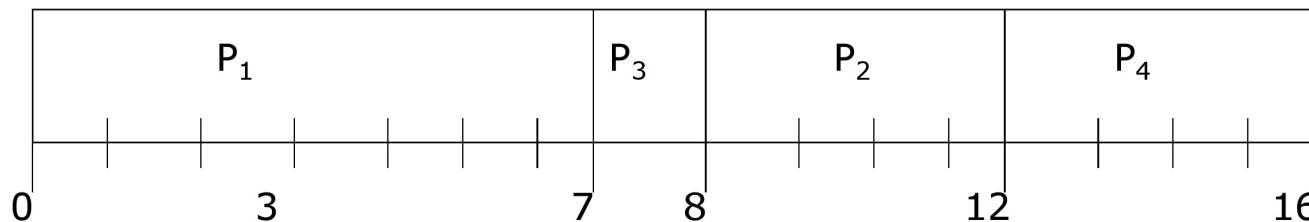
Avg. = $\boxed{7} \rightarrow$



Example of Non-Preemptive SJF

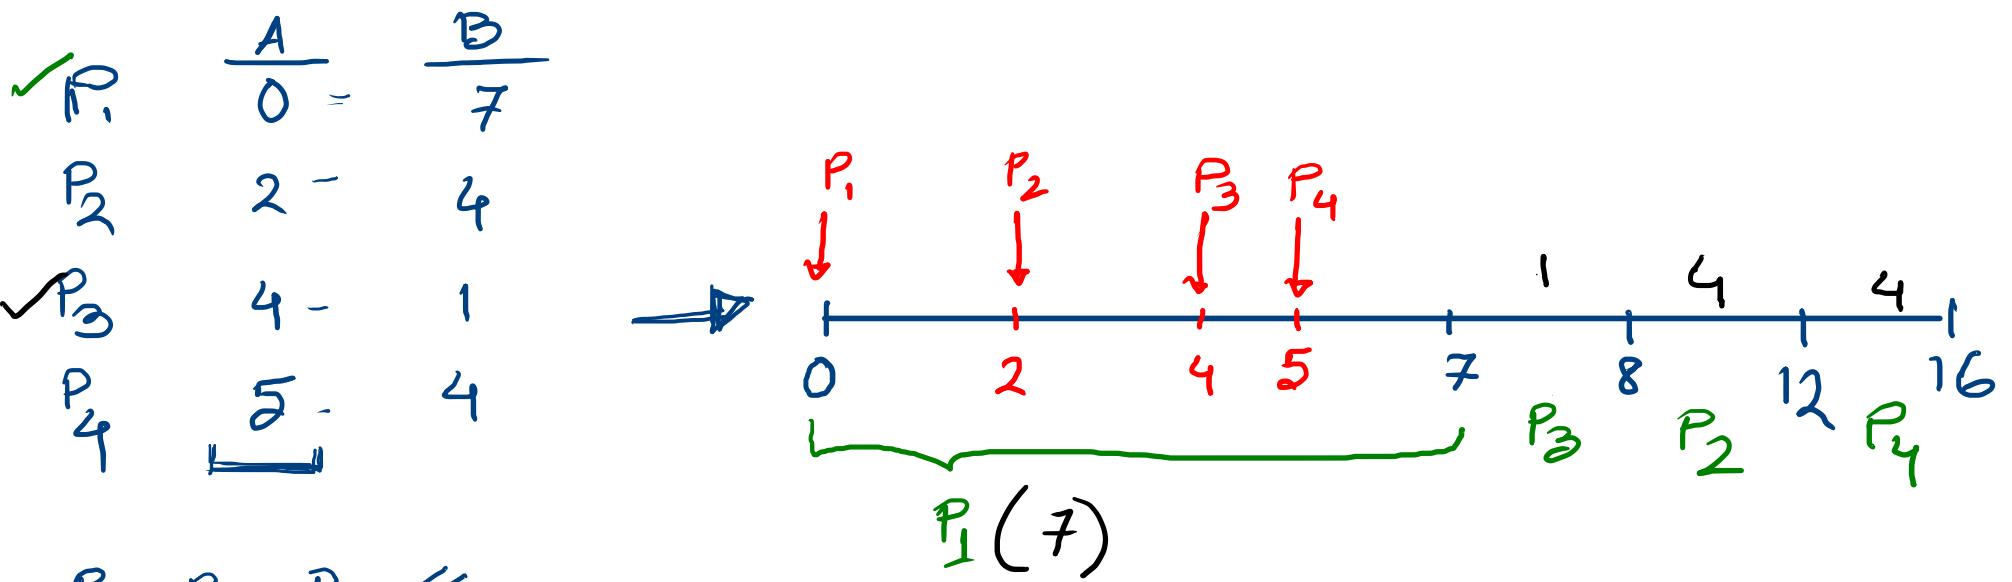
Process	Arrival Time	//	Burst Time
P_1	0.0		7
P_2	2.0		4
P_3	4.0		1
P_4	5.0		4

- SJF (non-preemptive)



$$\text{Average waiting time} = (0 + 6 + 3 + 7)/4 = 4$$





$P_3 \rightarrow P_2 \rightarrow P_4$

Waiting times.

$$P_1: 0 - 0 = 0$$

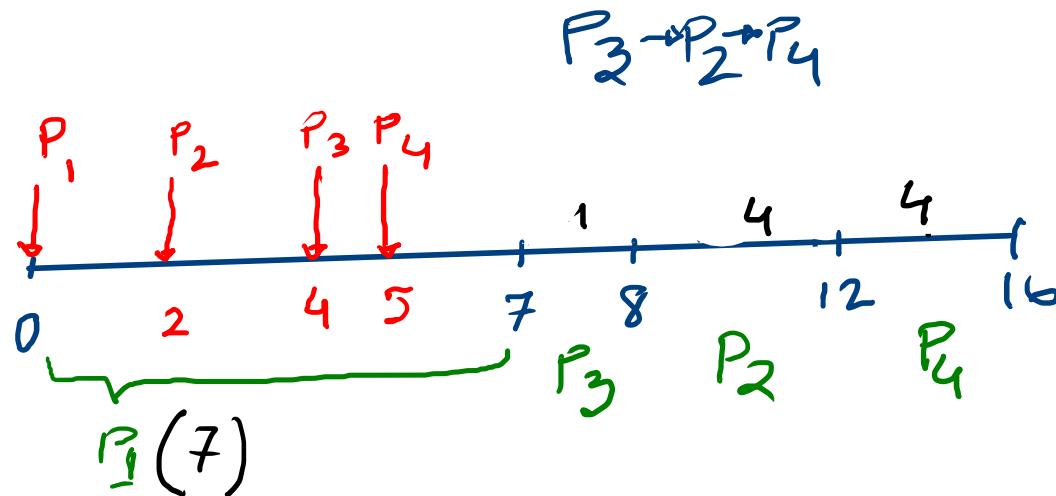
$$P_2: 8 - 2 = 6$$

$$P_3: 7 - 4 = 3$$

$$P_4: 12 - 5 = 7$$

$$\text{Avg: } 16/4 = 4$$

	A	B
P ₁	0	7
P ₂	2	4
P ₃	4	1
P ₄	5	4



Waiting times.

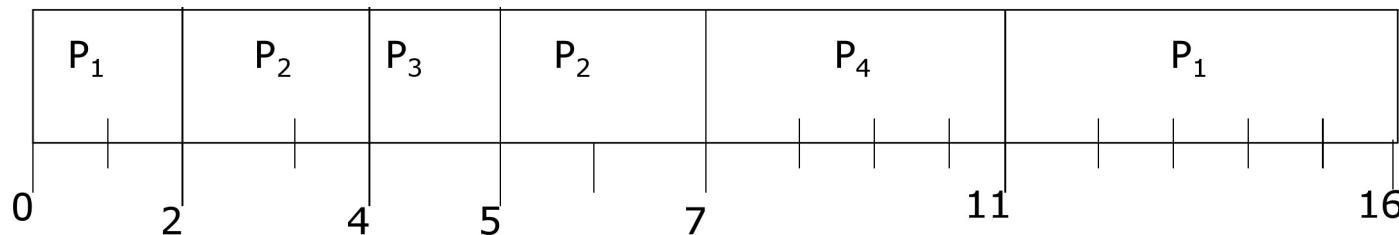
$$\left. \begin{array}{l} P_1 : 0 - 0 = 0 \\ P_2 : 8 - 2 = 6 \\ P_3 : 7 - 4 = 3 \\ P_4 : 12 - 5 = 7 \end{array} \right\} \text{Avg } = \underline{\underline{4}}$$



Example of Preemptive SJF

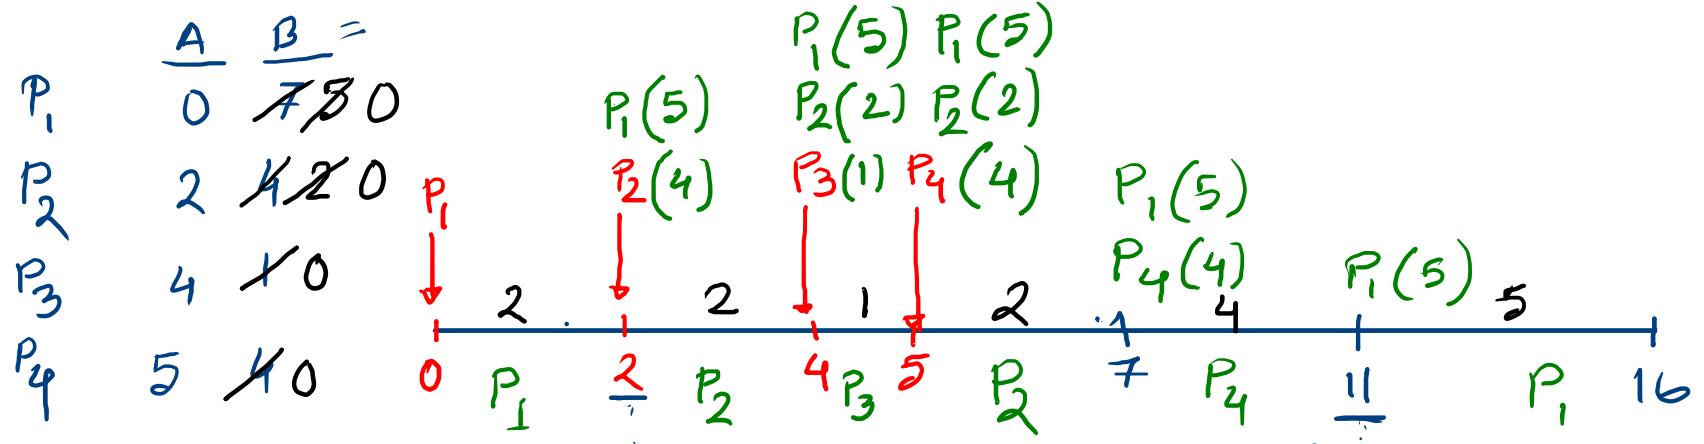
Process	Arrival Time	Burst Time
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

- SJF (preemptive)



- Average waiting time = $(9 + 1 + 0 + 2)/4 = \underline{3}$





Waiting times.

$$P_1: (0-0) + (11-2) = 9$$

$$P_2: (2-2) + (5-4) = 1$$

$$P_3: (4-4) = 0$$

$$P_4: (7-5) = 2$$

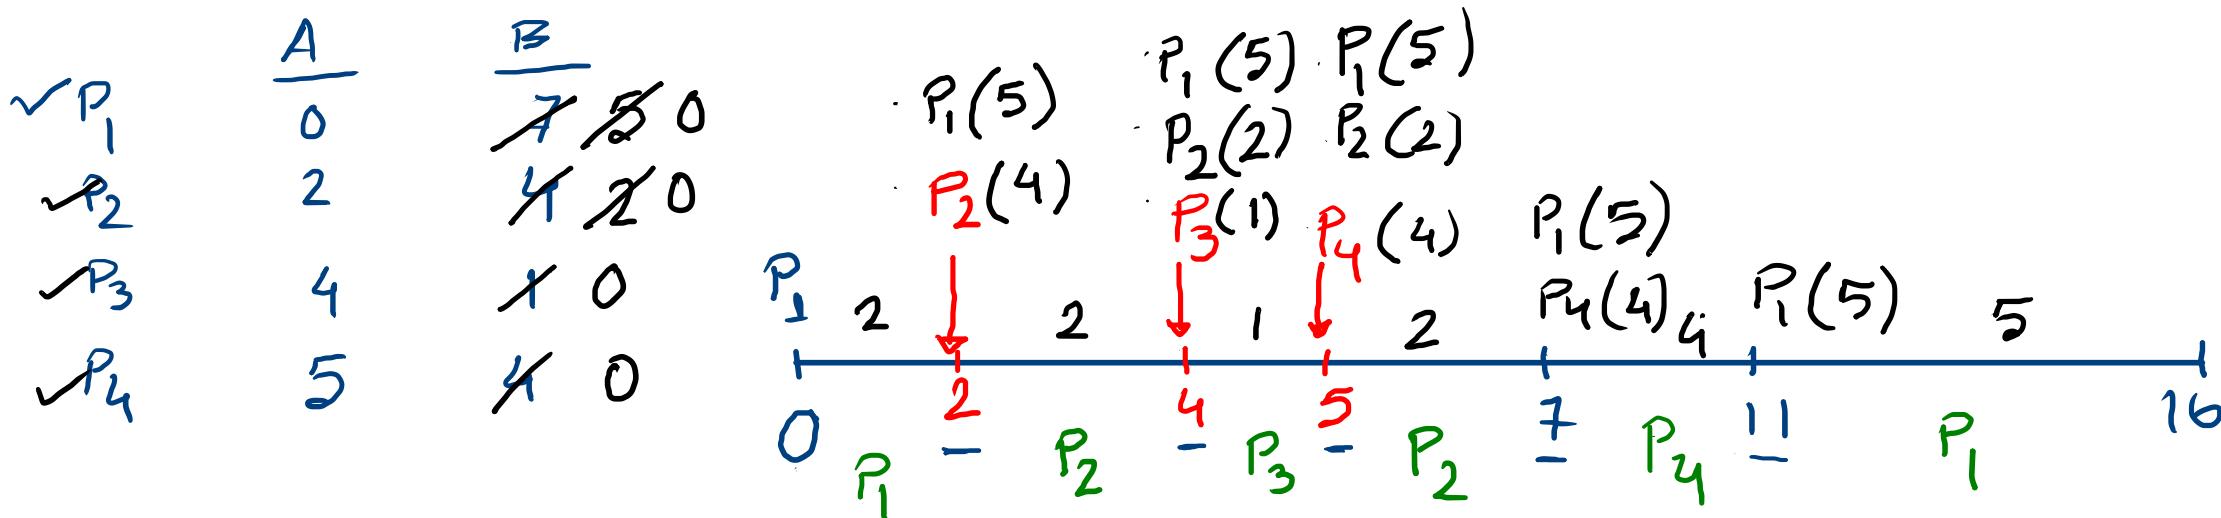
Avg: $\frac{9+1+0+2}{4} = 3$

$$\omega_t = (C_t - A_t) + \sum (CPU_A - CPU_{LR})$$

$$R_t = (C_t - A_t)$$

$$\text{Avg } \omega_t = \frac{1}{n} \sum [(C_{t_i} - A_{t_i}) + \sum (CPU_{A_i} - CPU_{LR_i})]$$

$$\text{Avg } R_t = \frac{1}{n} \sum (C_{t_i} - A_{t_i})$$



Waiting time:

$$P_1 = (0 - 0) + (11 - 2) = 9$$

$$P_2 = (2 - 2) + (5 - 4) = 1$$

$$P_3 = (4 - 4) = 0$$

$$P_4 = (7 - 5) = 2$$

Avg: $12/4 = 3$

* $\bar{W}_t = (C_t - A_t) + \sum (CPU_A - CPU_{LR})$

* $R_t = (C_t - A_t)$

* Avg $\bar{W}_t = \frac{1}{n} \sum [(C_{t_i} - A_{t_i}) + \sum (CPU_{A_i} - CPU_{LR_i})]$

* Avg $\bar{R}_t = \frac{1}{n} \sum (C_{t_i} - A_{t_i})$

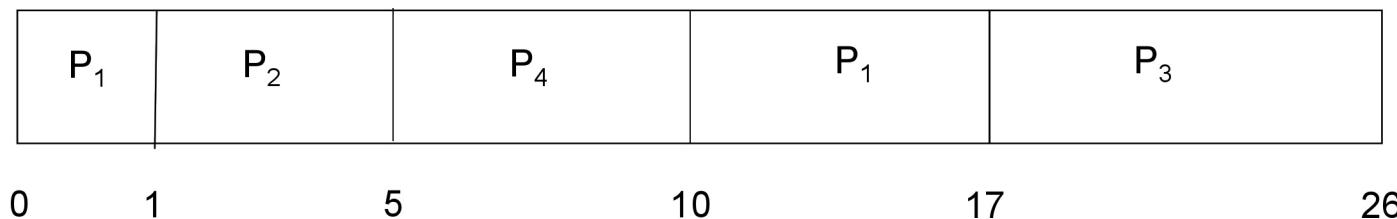


Example of Shortest-remaining-time-first

- Now we add the concepts of varying arrival times and preemption to the analysis

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

- Preemptive SJF Gantt Chart



- Average waiting time = $[(10-1)+(1-1)+(17-2)+5-3]/4 = 26/4 = 6.5 \text{ msec}$





Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer = highest priority)
 - ↗ Preemptive →
 - ▶ Will preempt the CPU if the priority of the newly arrived process is higher than that of the currently running process.
 - ↗ Nonpreemptive
 - ▶ Will simply put the new process at the head of the ready queue.
- ↗ SJF is priority scheduling where priority is the inverse of predicted next CPU burst time
- ↗ Priority can be defined either internally or externally.
 - ↗ Factors for internal priority assignment:
 - ▶ Time limit, memory requirements, the number of open files etc.
 - ↗ Factors for external priority assignment:
 - ▶ Importance of the process, the type and amount of funds being paid for computer use, department sponsoring works etc.



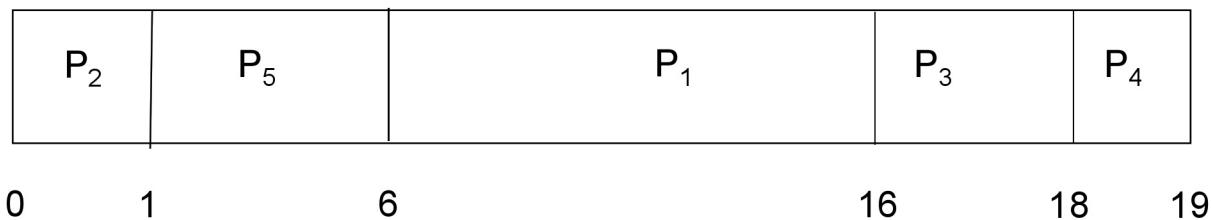


Example of Priority Scheduling

//

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2

- Priority scheduling Gantt Chart



Average waiting time = 8.2 msec





Priority Scheduling

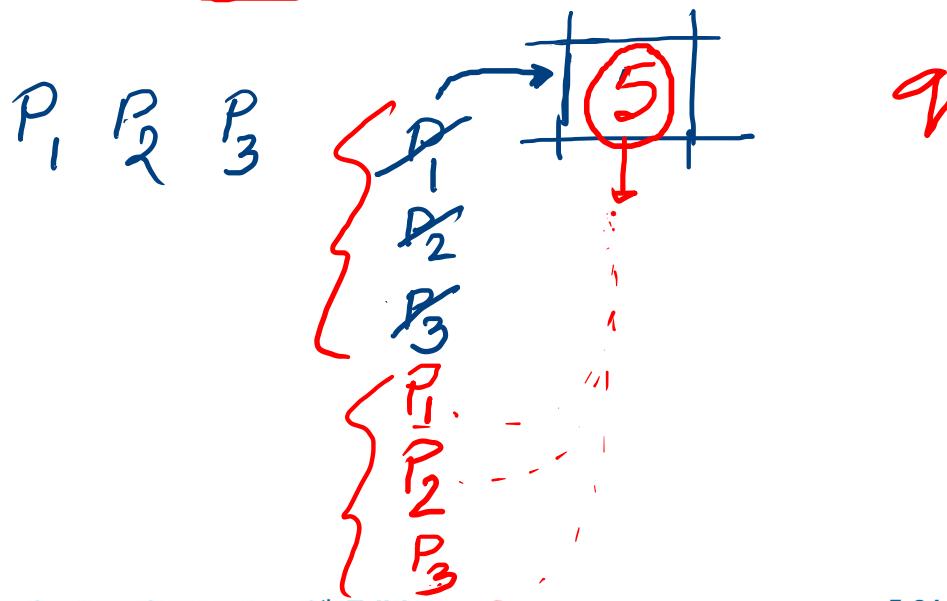
- Problem ≡ Starvation – low priority processes may never execute
- Solution ≡ Aging – as time progresses increase the priority of the process





Round Robin (RR)

- Each process gets a small unit of CPU time (**time quantum** q), usually 10-100 milliseconds.
- After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q , then
 - Each process gets $1/n$ of the CPU time in chunks of at most q time units at once.
 - No process waits more than $[(n-1) \times q]$ time units.
- Timer interrupts every quantum to schedule next process
- Performance
 - q large \Rightarrow FIFO
 - q small \Rightarrow q must be large with respect to context switch, otherwise overhead is too high

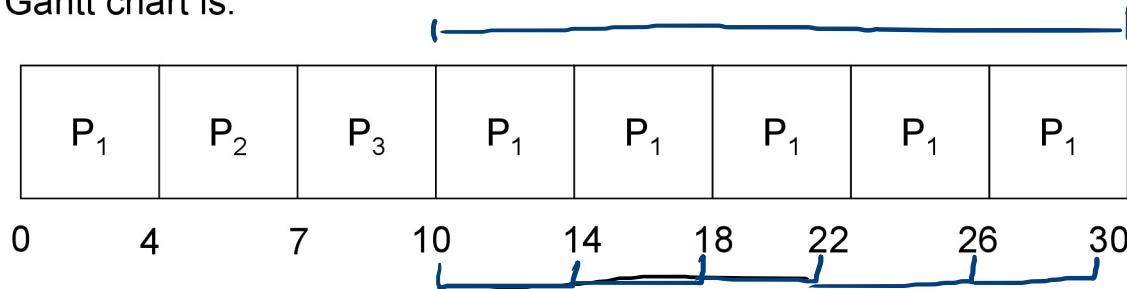




Example of RR with Time Quantum = 4

Process	Burst Time
P_1	24
P_2	3
P_3	3

- The Gantt chart is:

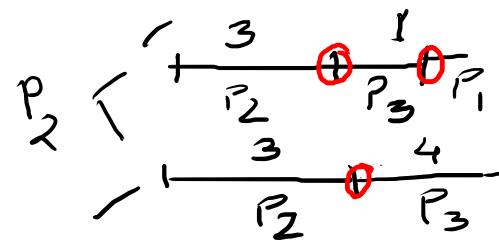
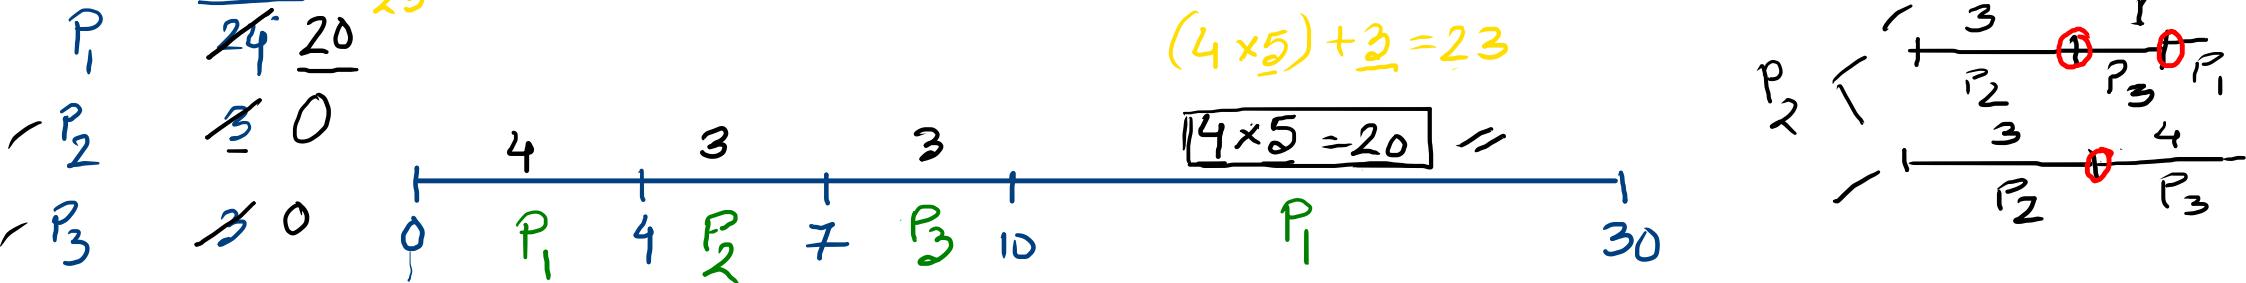


- Average waiting time is $17 / 3 = \underline{5.66}$ milisecond
- Typically, higher average turnaround than SJF, but better *response*
- quantum should be large compared to context switch time
- Quantum usually 10ms to 100ms, context switch < 10 usec



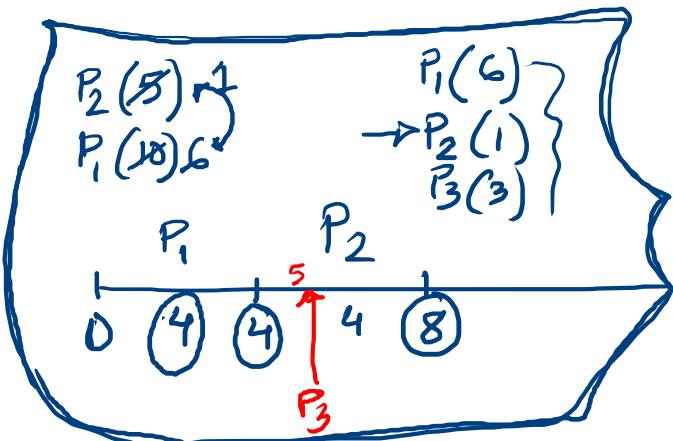
$$\frac{B}{24} \cdot 20$$

23



$$q = 4$$

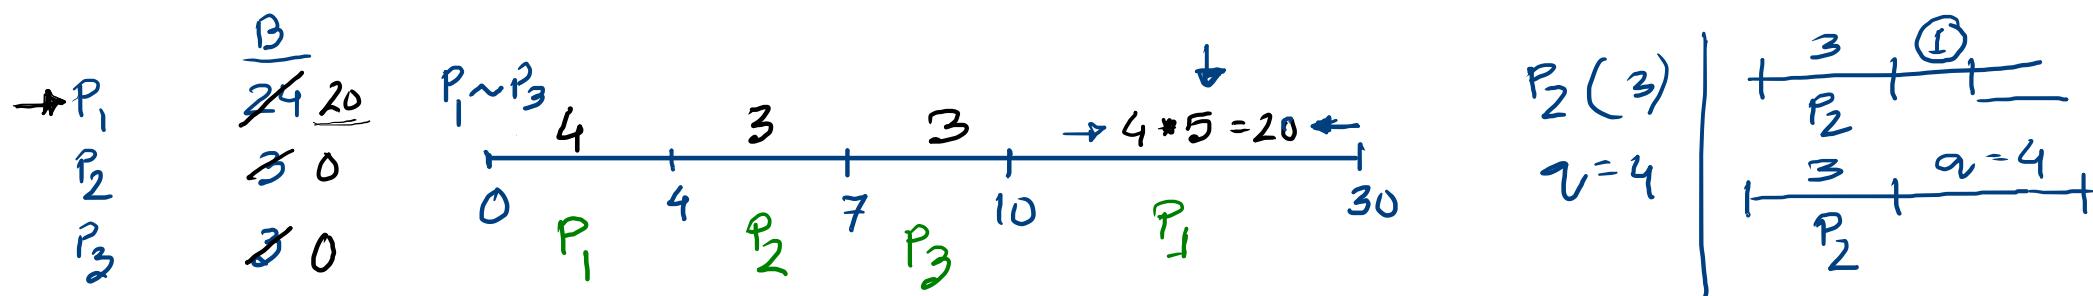
$$P_1 \cdot P_2 \cdot P_3$$



Waiting times.

$$\begin{aligned} P_1 &: (0-0) + (10-4) = 6 \\ P_2 &: (4-0) = 4 \\ P_3 &: (7-0) = 7 \end{aligned}$$

Avg: $\frac{7}{3} = 5.66$



$$\begin{array}{c}
 P_2(3) \\
 v=4 \\
 \hline
 +\frac{3}{P_2} + \textcircled{1} \\
 \hline
 3 \\
 \hline
 P_2
 \end{array}$$

$$q = 4$$

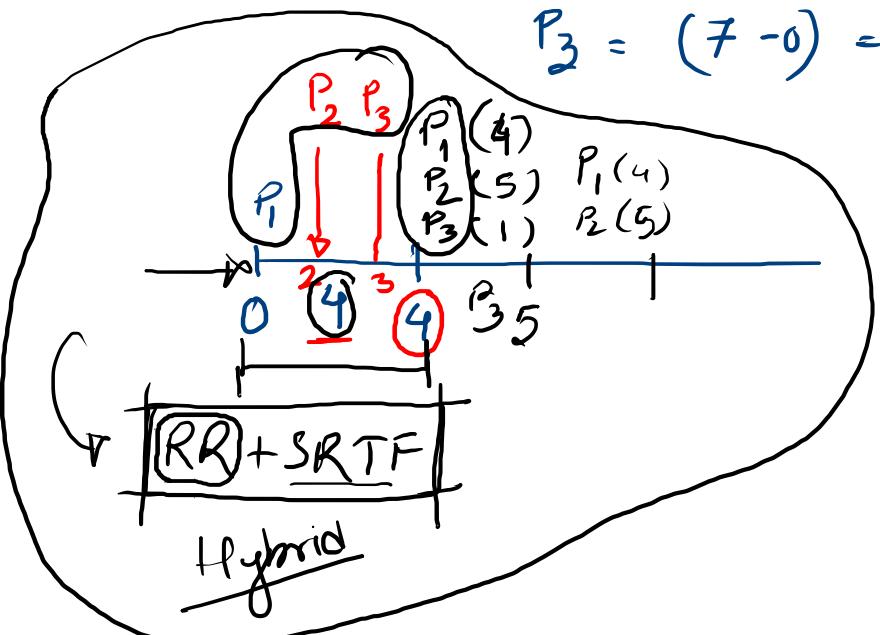
Waiting time

$$P_1 = (0 - 0) + (10 - 4) = 6$$

$$P_2 = (4 - 0) = 4$$

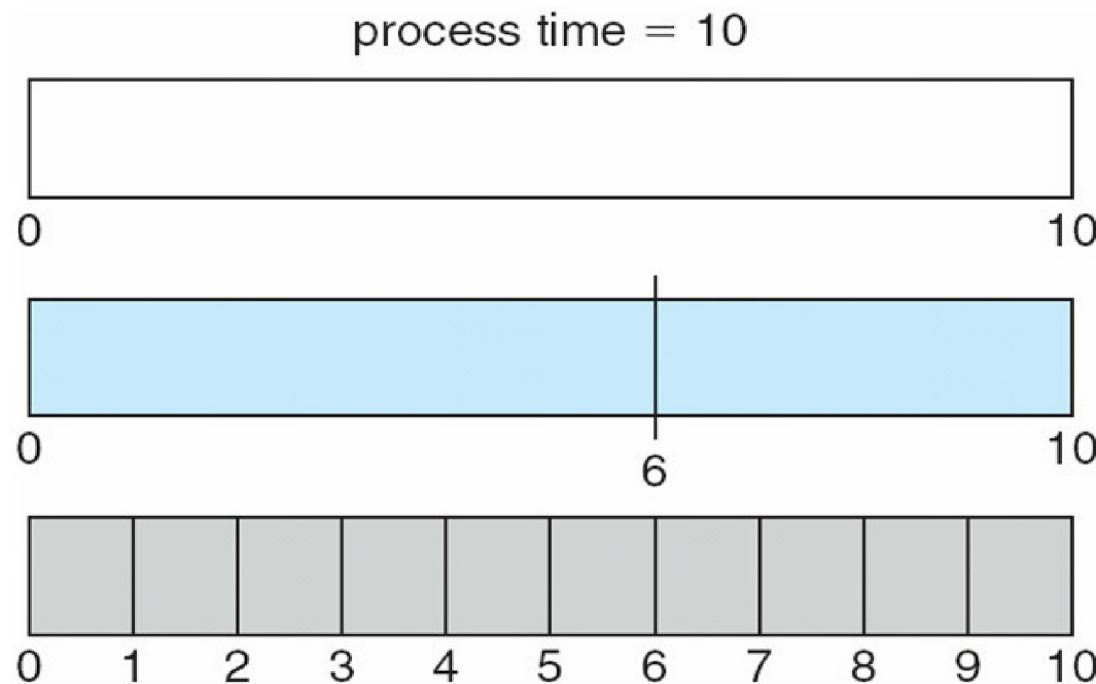
$$P_3 = (7 - 0) = 7$$

Avg: $\frac{17}{3} = 5.66$





Time Quantum and Context Switch Time



quantum

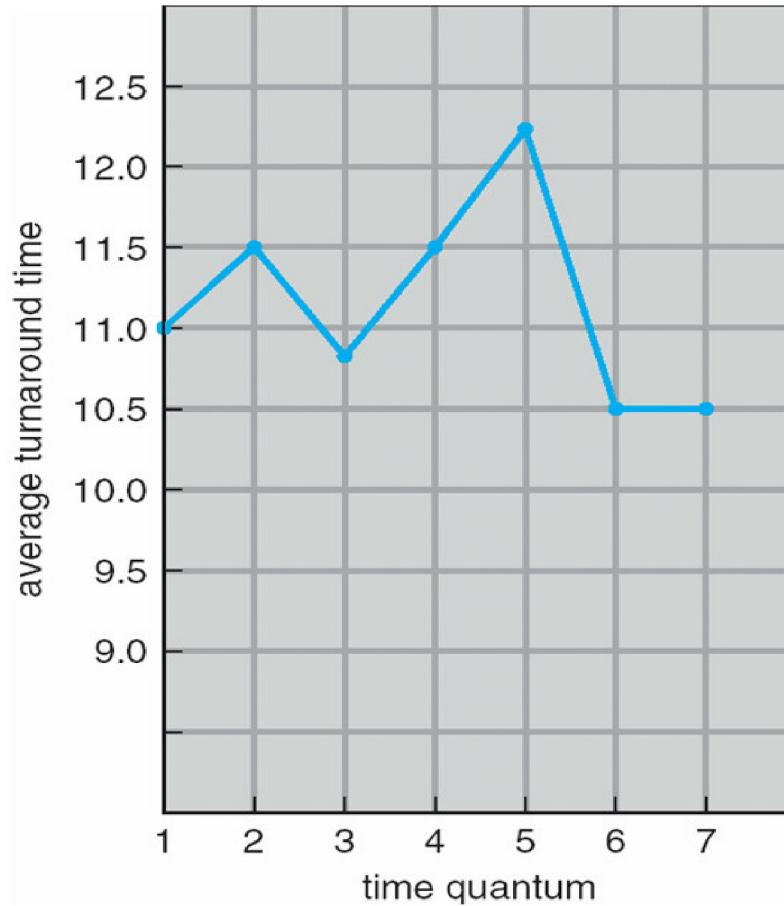


context switches





Turnaround Time Varies With The Time Quantum



process	time
P_1	6
P_2	3
P_3	1
P_4	7

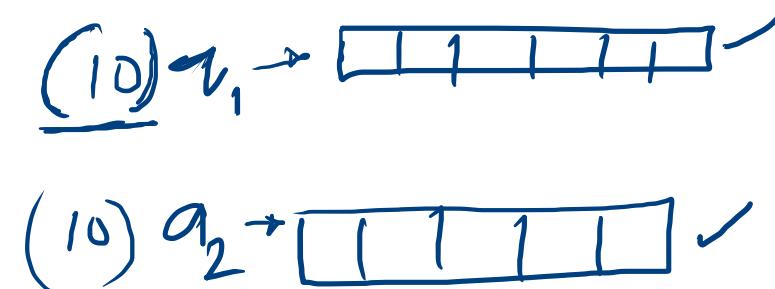
80% of CPU bursts should be shorter than quantum





Multilevel Queue

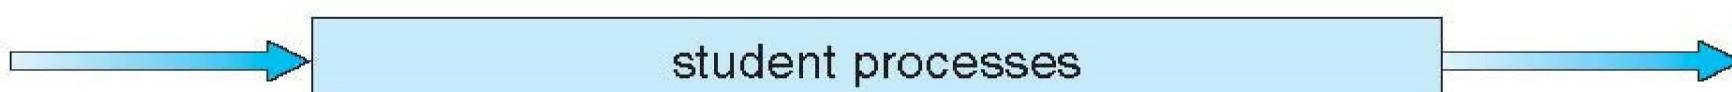
- Another class of scheduling algorithm needs- in which processes are classified into different groups, e.g.:
 - foreground (interactive) processes
 - background (batch) processes
- They have different response time requirements-so different scheduling needs.
- Foreground processes may have priority over background processes.
- A multilevel queue-scheduling algorithm partitions the ready queue into several separate queues-we can see it in the figure of next slide:-
- → Each queue has its own scheduling algorithm:
 - Foreground queue scheduled by – RR algorithm
 - Background queue scheduled by – FCFS algorithm
- Scheduling must be done between the queues:
 - Fixed priority preemptive scheduling; (i.e., serve all from foreground then from background). Possibility of starvation.
 - Time slice– each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., foreground queue can be given 80% of the CPU time for RR-scheduling among its processes, while 20% to background in FCFS manner.





Multilevel Queue Scheduling

highest priority



lowest priority

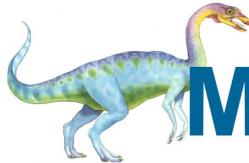




Multilevel Feedback Queue scheduling

- In multilevel queue, processes do not move from one queue to the other----But
- Multilevel Feedback Queue scheduling, allows a process to move between queues.
- If a process uses too much CPU time, it will be moved to a lower priority queue.
- Similarly, a process that waits too long in a lower-priority queue may me moved to a higher-priority queue.





Multilevel Feedback Queue scheduling

- Multilevel-feedback-queue scheduler defined by the following parameters:

- Multilevel-feedback-queue scheduler defined by the following parameters:
 - number of queues
 - scheduling algorithms for each queue
 - method used to determine when to upgrade a process
 - method used to determine when to demote a process
 - method used to determine which queue a process will enter when that process needs service





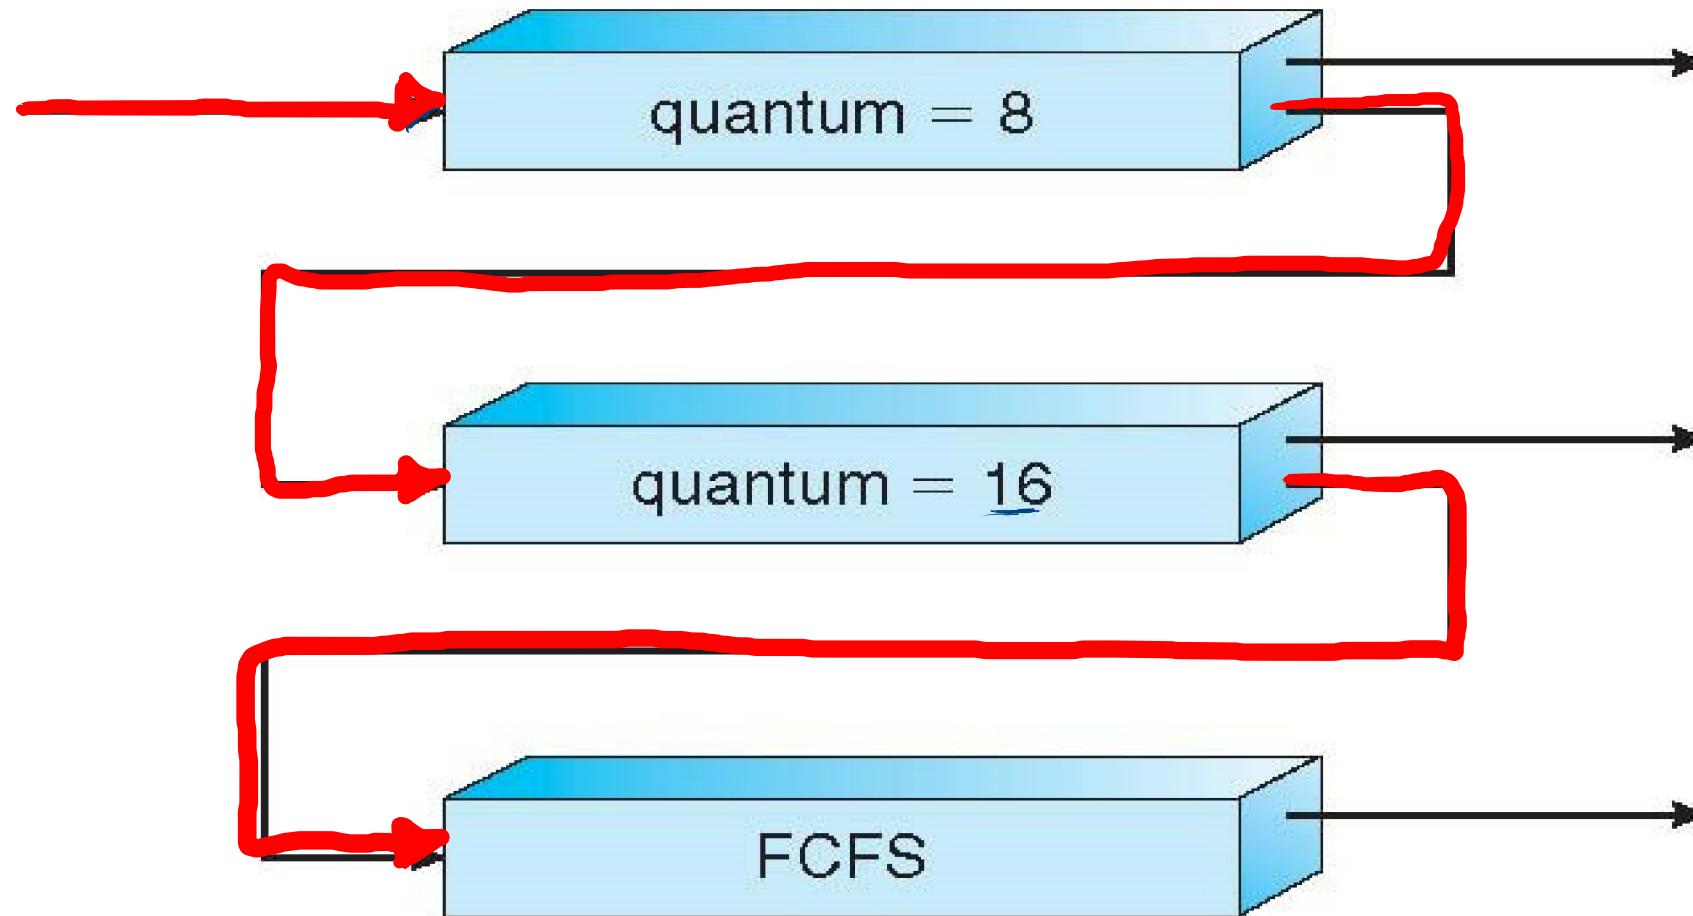
Example of Multilevel Feedback Queue

- Three queues: (can see the figure in next slide)
 - Q_0 – RR with time quantum 8 milliseconds
 - Q_1 – RR time quantum 16 milliseconds
 - Q_2 – FCFS
- Scheduling
 - A new job enters queue Q_0 which is served for RR
 - ▶ When it gains CPU, job receives 8 milliseconds
 - ▶ If it does not finish in 8 milliseconds, job is moved to queue Q_1
 - At Q_1 job is again served RR and receives 16 additional milliseconds
 - ▶ If it still does not complete, it is preempted and moved to queue Q_2





Multilevel Feedback Queues



End of Chapter 6

